

Recommendation: Major revision before scientific publication.

The manuscript presents an ambitious modular cognitive architecture combining multi-agent reasoning, persistent “cocoon” memory, ethical evaluation, and meta-cognitive strategy generation. The paper’s strongest aspects are its clear architectural framing, explicit ablation-style conditions, inclusion of paired statistical analysis, and unusually candid limitations section. The work reports a large composite-score improvement for full Codette over a single-agent baseline on a 17-problem benchmark, while also acknowledging that memory and strategy-evolution increments are not statistically significant at the present sample size.

However, the current evidence is not yet strong enough to support the broader scientific claims. The main concerns are construct validity of the automated scoring, small and purpose-built benchmark size, weak baseline comparisons, inconsistent reported numbers across sections, and a theoretical framework that is mostly metaphorical unless stronger implementation-level assumptions are demonstrated. The manuscript is promising as a systems/workshop paper, but it needs substantial tightening before it can support claims of generalizable reasoning improvement.

Strengths:

1. The four-condition comparison—SINGLE, MULTI, MEMORY, and CODETTE—is useful and interpretable. The layered architecture and agent roles are described in enough detail to understand the intended system flow.
2. The paper avoids claiming a general convergence guarantee and instead frames the $RC+\xi$ formalism as conditional on assumptions such as bounded domains, projection, and contraction. This restraint improves credibility.
3. The manuscript states that MEMORY vs. MULTI and CODETTE vs. MEMORY do not reach significance after Holm correction, which is important because those are central claimed innovations.
4. The paper explicitly acknowledges automated scoring, template-based agents, single-model evaluation, novelty measurement, and the lack of human evaluation as limitations. This section is scientifically valuable and should be preserved.
5. The radar plot on page 9, reasoning-depth plot on page 10, latency-performance plot on page 12, and benchmark bar chart on page 18 help readers quickly understand the reported trends. That said, the architecture diagram on page 5 is too minimal to carry much explanatory weight.

Major comments

1. The paper reports a large improvement from SINGLE to CODETTE, but the benchmark contains only 17 problems and uses automated scoring. This makes the result useful as an internal diagnostic, but not yet strong evidence of general reasoning improvement. The paper itself notes that automated metrics may not correlate with expert judgment or downstream utility.
2. The main results report $SINGLE = 0.356$ and $CODETTE = 0.689$, while the appendix table reports $SINGLE = 0.338$ and $CODETTE = 0.652$. Similarly, the memory impact is reported as $+0.018$ in the main analysis but later described as $+0.004$ in the limitations. These inconsistencies create uncertainty about which run is authoritative.

3. The SINGLE baseline is a single analytical agent without memory or synthesis. This is informative as a component ablation, but it is not a strong external baseline. The paper should compare Codette against stronger alternatives such as standard chain-of-thought prompting, self-consistency, Reflexion-style self-critique, debate-style multi-agent prompting, RAG/memory augmentation without strategy synthesis, and a single-agent response with matched token budget.
4. The scoring dimensions include reasoning depth, perspective diversity, novelty, factual grounding, and Turing naturalness. These are reasonable dimensions, but the manuscript does not sufficiently demonstrate that the scorer measures true reasoning quality rather than response length, keyword density, or template structure. The output-size table shows that CODETTE responses are much longer than SINGLE responses, which is a serious confound.
5. The RC+ ξ model is mathematically framed, but the paper acknowledges that ethical checks are implemented as discrete rubric/heuristic gates rather than differentiable gradients, and that convergence requires assumptions not yet established for the real system.
6. The most novel parts of the system are cocoon memory and meta-cognitive strategy evolution, but the paired results show that MEMORY vs. MULTI and CODETTE vs. MEMORY are not statistically significant.
7. The system uses Llama 3.1 8B, LoRA adapters, SQLite/FTS5 memory, routing logic, heuristic scoring, and strategy generation, but many details needed for replication are missing or deferred to external artifacts.

Minor comments

1. The title is long and may overstate “convergent” unless convergence is formally demonstrated. Consider: **“Codette: A Multi-Perspective Cognitive Architecture with Memory and Meta-Cognitive Strategy Evolution.”**
2. The term **“quantum-probabilistic”** should be justified technically or renamed to avoid sounding metaphorical.
3. The architecture figure on page 5 is too sparse. Replace it with a diagram showing query input, routing, agents, memory retrieval, AEGIS checks, critic/synthesis, strategy generation, and output validation.
4. Define whether agent weights are fixed, learned, heuristic, or classifier-derived.
5. The AEGIS ethics framework is interesting, but the manuscript should avoid implying cultural universality. The paper already notes this issue; the final version should either ground the Ubuntu/sustainability components in scholarship or narrow the claim.
6. The paper should distinguish more carefully between **observed benchmark score gains**, **reasoning-quality gains**, and **real-world utility gains**.
7. Some formatting artifacts appear in the PDF text, such as broken words and unusual hyphenation. These should be cleaned before submission.

Final recommendation

Major revision. The manuscript has a promising architecture and a useful experimental scaffold, but the current scientific evidence is preliminary. The paper will be much stronger if it narrows its claims, resolves numerical inconsistencies, validates the scoring system with humans, adds stronger

baselines, controls for response length, and clarifies the gap between the dynamical-system formalism and the implemented software.